

Kevin Berry

(678) 237-3418 · Alpharetta, GA

kpberry11@gmail.com · [linkedin](#) · [github](#) · kpberry.github.io

EDUCATION

Georgia Institute of Technology, B.S. Computer Science

Aug 2015 - Dec 2018

Selected Coursework: Machine Learning, Linguistics, Computer Vision, Compilers, Linear Algebra

Extracurriculars: Communications officer for The Agency (ML/AI club) and Big O (Theoretical CS club)

GPA: 3.96/4.0

SKILLS

Programming Languages: Python, Rust, JavaScript, C

Tools: PyTorch, pandas, networkx, scikit-learn, Docker, Git, Azure, SQL, MongoDB, \LaTeX

Human Languages: English (fluent), Portuguese (intermediate), Spanish (intermediate)

EXPERIENCE

Machine Learning Engineer

Jun 2017 - Present

Worthix

Alpharetta, GA

Here, my goal was to help companies understand customer feedback. I have been responsible for designing, developing, deploying, monitoring, and maintaining a wide variety of tools to accomplish this goal. I recruited and currently direct a small team which works on these tasks with me. The following are some highlights of my individual contributions:

- Developed LLM-based API for large scale multi-level feedback analysis and summarization
- Developed and trained NLP models and API for real-time survey response topic classification
- Developed novel algorithms for large-scale unsupervised graph classification
- Developed queries and statistical analyses for client-facing dashboards
- Deployed and monitored APIs and applications as scalable microservices in Azure
- Conducted scientific experiments to improve survey data and model quality
- Developed fast Monte-Carlo Shapley value approximation for calculating response topic importance
- Developed reverse-geolocation API capable of resolving millions of point-in-region queries per second
- Contributed Python code to networkx which massively improved performance for graph unions and intersections
- Contributed Rust code to polars dataframe library which solved infinite loop in core groupby operations

Teaching Assistant

May 2016 - Dec 2017

Computer Organization and Programming (CS 2110)

Atlanta, GA

- Led recitations on C, Assembly, CPU datapaths, and digital logic
- Wrote software to automate grading of Java programs and circuits

PROJECTS

image-to-ascii (Rust): Uses computer vision techniques to convert images and gifs to ASCII art. Capable of converting 120+ images per second. Uses hand-coded SIMD instructions for improved performance. ([Github](#)) ([Crates.io](#)) ([Github](#))

Rolling a d_7 in Finite(ish) Time: Paper analyzing the problem of simulating dice with any number of faces given only a standard die set. Proves some efficiency results for interesting cases. ([Paper](#))

detect-duplicates (Rust): Recursively finds all of the duplicated files in a directory. Reads each file at most twice and, in expectation, only has one file loaded at a time. Can check the entire Linux kernel repository (roughly 5GB, over 70,000 files) in about 5 seconds. ([Github](#)) ([Crates.io](#))

Readability Analyzer (JavaScript): Assigns average text grade level based on standard readability metrics. Uses MLP with hand-engineered features to count syllables. Can analyze 2,000,000+ characters per second. ([Web App](#))

Vertex Cover Algorithms (Python): Implemented algorithms to solve minimum vertex cover problem. Made a new approximation algorithm which achieved state of the art accuracy on 10 out of 11 real-world datasets. ([Paper](#))